

```
1 !pip install gymnasium
```

```
Requirement already satisfied: gymnasium in /usr/local/lib/python3.12/dist-packages (1.2.0)
Requirement already satisfied: numpy>=1.21.0 in /usr/local/lib/python3.12/dist-packages (from gymnasium) (2.0.2)
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from gymnasium) (3.1.1)
Requirement already satisfied: typing-extensions>=4.3.0 in /usr/local/lib/python3.12/dist-packages (from gymnasium) (4.15.0)
Requirement already satisfied: farama-notifications>=0.0.1 in /usr/local/lib/python3.12/dist-packages (from gymnasium) (0.0.4)
```

```
1 pip install git+https://github.com/mimoralea/gym-walk#egg=gym-walk
```

```
Collecting gym-walk
  Cloning https://github.com/mimoralea/gym-walk to /tmp/pip-install-saudrrye/gym-walk_88175d5cab746a8b8dabd0b7d18f915
  Running command git clone --filter=blob:none --quiet https://github.com/mimoralea/gym-walk /tmp/pip-install-saudrrye/gym-walk
  Resolved https://github.com/mimoralea/gym-walk to commit b915b94cf2ad16f8833a1ad92ea94e88159279f5
  Preparing metadata (setup.py) ... done
Requirement already satisfied: gym in /usr/local/lib/python3.12/dist-packages (from gym-walk) (0.25.2)
Requirement already satisfied: numpy>=1.18.0 in /usr/local/lib/python3.12/dist-packages (from gym->gym-walk) (2.0.2)
Requirement already satisfied: cloudpickle>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from gym->gym-walk) (3.1.1)
Requirement already satisfied: gym-notices>=0.0.4 in /usr/local/lib/python3.12/dist-packages (from gym->gym-walk) (0.1.0)
Building wheels for collected packages: gym-walk
  Building wheel for gym-walk (setup.py) ... done
  Created wheel for gym-walk: filename=gym_walk-0.0.2-py3-none-any.whl size=5377 sha256=a9ca47f2bce0f9f7dc1ef3fec65a06e404fac97
  Stored in directory: /tmp/pip-ephem-wheel-cache-4p4ra7on/wheels/bf/23/e5/a94be4a90dd18f7ce958c21f192276cb01ef0daaf2bc66583b
Successfully built gym-walk
Installing collected packages: gym-walk
Successfully installed gym-walk-0.0.2
```

```
1 import warnings ; warnings.filterwarnings('ignore')
2
3 import gym, gym_walk
4 import numpy as np
5
6 import random
7 import warnings
8
9 warnings.filterwarnings('ignore', category=DeprecationWarning)
10 np.set_printoptions(suppress=True)
11 random.seed(123); np.random.seed(123)
12
```

Gym has been unmaintained since 2022 and does not support NumPy 2.0 amongst other critical functionality. Please upgrade to Gymnasium, the maintained drop-in replacement of Gym, or contact the authors of your software and request the See the migration guide at https://gymnasium.farama.org/introduction/migration_guide/ for additional information.

```
1 def print_policy(pi, P, action_symbols=('<', 'v', '>', '^'), n_cols=4, title='Policy:'):
2     print(title)
3     arrs = {k:v for k,v in enumerate(action_symbols)}
4     for s in range(len(P)):
5         a = pi(s)
6         print("| ", end="")
7         if np.all([done for action in P[s].values() for _, _, done in action]):
8             print("".rjust(9), end=" ")
9         else:
10            print(str(s).zfill(2), arrs[a].rjust(6), end=" ")
11            if (s + 1) % n_cols == 0: print("|")
```

```
1 def print_state_value_function(V, P, n_cols=4, prec=3, title='State-value function:'):
2     print(title)
3     for s in range(len(P)):
4         v = V[s]
5         print("| ", end="")
6         if np.all([done for action in P[s].values() for _, _, done in action]):
7             print("".rjust(9), end=" ")
8         else:
9             print(str(s).zfill(2), '{}'.format(np.round(v, prec)).rjust(6), end=" ")
10            if (s + 1) % n_cols == 0: print("|")
```

```
1 def probability_success(env, pi, goal_state, n_episodes=100, max_steps=200):
2     random.seed(123); np.random.seed(123) ; env.seed(123)
3     results = []
4     for _ in range(n_episodes):
5         state, done, steps = env.reset(), False, 0
6         while not done and steps < max_steps:
7             state, _, done, h = env.step(pi(state))
8             steps += 1
9         results.append(state == goal_state)
10    return np.sum(results)/len(results)
```

```

1 def mean_return(env, pi, n_episodes=100, max_steps=200):
2     random.seed(123); np.random.seed(123) ; env.seed(123)
3     results = []
4     for _ in range(n_episodes):
5         state, done, steps = env.reset(), False, 0
6         results.append(0.0)
7         while not done and steps < max_steps:
8             state, reward, done, _ = env.step(pi(state))
9             results[-1] += reward
10            steps += 1
11    return np.mean(results)

```

```

1 env = gym.make('FrozenLake-v1')
2 P = env.env.P
3 init_state = env.reset()
4 goal_state = 15
5 LEFT, DOWN, RIGHT, UP = range(4)
6

```

```

1 P
(0.3333333333333333, 13, 0.0, False)],
1: [(0.3333333333333333, 8, 0.0, False),
(0.3333333333333333, 13, 0.0, False),
(0.3333333333333333, 10, 0.0, False)],
2: [(0.3333333333333333, 13, 0.0, False),
(0.3333333333333333, 10, 0.0, False),
(0.3333333333333333, 5, 0.0, True)],
3: [(0.3333333333333333, 10, 0.0, False),
(0.3333333333333333, 5, 0.0, True),
(0.3333333333333333, 8, 0.0, False)],
10: {0: [(0.3333333333333333, 6, 0.0, False),
(0.3333333333333333, 9, 0.0, False),
(0.3333333333333333, 14, 0.0, False)],
1: [(0.3333333333333333, 9, 0.0, False),
(0.3333333333333333, 14, 0.0, False),
(0.3333333333333333, 11, 0.0, True)],
2: [(0.3333333333333333, 14, 0.0, False),
(0.3333333333333333, 11, 0.0, True),
(0.3333333333333333, 6, 0.0, False)],
3: [(0.3333333333333333, 11, 0.0, True),
(0.3333333333333333, 6, 0.0, False),
(0.3333333333333333, 9, 0.0, False)]},
11: {0: [(1.0, 11, 0, True)],
1: [(1.0, 11, 0, True)],
2: [(1.0, 11, 0, True)],
3: [(1.0, 11, 0, True)]},
12: {0: [(1.0, 12, 0, True)],
1: [(1.0, 12, 0, True)],
2: [(1.0, 12, 0, True)],
3: [(1.0, 12, 0, True)]},
13: {0: [(0.3333333333333333, 9, 0.0, False),
(0.3333333333333333, 12, 0.0, True),
(0.3333333333333333, 13, 0.0, False)],
1: [(0.3333333333333333, 12, 0.0, True),
(0.3333333333333333, 13, 0.0, False),
(0.3333333333333333, 14, 0.0, False)],
2: [(0.3333333333333333, 13, 0.0, False),
(0.3333333333333333, 14, 0.0, False),
(0.3333333333333333, 9, 0.0, False)],
3: [(0.3333333333333333, 14, 0.0, False),
(0.3333333333333333, 9, 0.0, False),
(0.3333333333333333, 12, 0.0, True)]},
14: {0: [(0.3333333333333333, 10, 0.0, False),
(0.3333333333333333, 13, 0.0, False),
(0.3333333333333333, 14, 0.0, False)],
1: [(0.3333333333333333, 13, 0.0, False),
(0.3333333333333333, 14, 0.0, False),
(0.3333333333333333, 15, 1.0, True)],
2: [(0.3333333333333333, 14, 0.0, False),
(0.3333333333333333, 15, 1.0, True),
(0.3333333333333333, 10, 0.0, False)],
3: [(0.3333333333333333, 15, 1.0, True),
(0.3333333333333333, 10, 0.0, False),
(0.3333333333333333, 13, 0.0, False)]},
15: {0: [(1.0, 15, 0, True)],
1: [(1.0, 15, 0, True)],
2: [(1.0, 15, 0, True)],
3: [(1.0, 15, 0, True)]}

```

```

1 init_state

```

```

0

```

```

1 state, reward, done, info = env.step(RIGHT)
2 print("state:{0} - reward:{1} - done:{2} - info:{3}".format(state, reward, done, info))

```

```
state:4 - reward:0.0 - done:False - info:{'prob': 0.3333333333333333}
```

```

1 pi_frozenlake1 = lambda s: {
2     0: RIGHT,
3     1: RIGHT,
4     2: RIGHT,
5     3: RIGHT,
6     4: RIGHT,
7     5: RIGHT,
8     6: RIGHT,
9     7: RIGHT,
10    8: RIGHT,
11    9: RIGHT,
12    10: RIGHT,
13    11: RIGHT,
14    12: RIGHT,
15    13: RIGHT,
16    14: RIGHT,
17    15: RIGHT #Stop
18 }[s]
19 print("Name: JERUSHLIN JOSE JB ")
20 print("Register Number: 21222240039 ")
21 print_policy(pi_frozenlake1, P, action_symbols=('<', 'v', '>', '^'), n_cols=4)

```

Name: JERUSHLIN JOSE JB

Register Number: 21222240039

Policy:

00	> 01	> 02	> 03	>
04	>	06	>	
08	> 09	> 10	>	
	13	> 14	>	

```

1 print('Reaches goal {:.2f}%. Obtains an average undiscounted return of {:.4f}'.format(
2     probability_success(env, pi_frozenlake1, goal_state=goal_state)*100,
3     mean_return(env, pi_frozenlake1)))

```

Reaches goal 1.00%. Obtains an average undiscounted return of 0.0100.

```

1 def policy_evaluation(pi, P, gamma=1.0, theta=1e-10):
2     prev_V = np.zeros(len(P), dtype=np.float64)
3     while True:
4         V = np.zeros(len(P), dtype=np.float64)
5         for s in range(len(P)):
6             for prob, next_state, reward, done in P[s][pi(s)]:
7                 V[s] += prob * (reward + gamma * prev_V[next_state] * (not done))
8             if np.max(np.abs(prev_V - V)) < theta:
9                 break
10        prev_V = V.copy()
11    return V

```

```

1 V1 = policy_evaluation(pi_frozenlake1, P)
2 print("Name: JERUSHLIN JOSE JB ")
3 print("Register Number: 21222240039 ")
4 print_state_value_function(V1, P, n_cols=4, prec=5)

```

Name: JERUSHLIN JOSE JB

Register Number: 21222240039

State-value function:

00 0.0315	01 0.02381	02 0.04762	03 0.0	
04 0.03919		06 0.09524		
08 0.08608	09 0.21905	10 0.2381		
	13 0.41905	14 0.61905		

```

1 def policy_improvement(V, P, gamma=1.0):
2     Q = np.zeros((len(P), len(P[0])), dtype=np.float64)
3     # Write your code here to improve the given policy
4     for s in range(len(P)):
5         for a in range(len(P[s])):
6             for prob, next_state, reward, done in P[s][a]:
7                 Q[s][a] += prob * (reward + gamma * V[next_state] * (not done))
8     new_pi = lambda s: {s:a for s, a in enumerate(np.argmax(Q, axis=1))}[s]
9     return new_pi

```

```

1 pi_2 = policy_improvement(V1, P)
2 print("Name: JERUSHLIN JOSE JB ")
3 print("Register Number: 21222240039 ")
4 print_policy(pi_2, P, action_symbols=('<', 'v', '>', '^'), n_cols=4)

```

Name: JERUSHLIN JOSE JB
Register Number: 21222240039

Policy:

00	<	01	^	02	<	03	<
04	<			06	<		
08	^	09	v	10	<		
		13	>	14	v		

```

1 print('Reaches goal {:.2f}%. Obtains an average undiscounted return of {:.4f}.'.format(
2     probability_success(env, pi_2, goal_state=goal_state)*100,
3     mean_return(env, pi_2)))

```

Reaches goal 66.00%. Obtains an average undiscounted return of 0.6600.

```

1 V2 = policy_evaluation(pi_2, P)
2 print("Name: JERUSHLIN JOSE JB ")
3 print("Register Number: 21222240039 ")
4 print_state_value_function(V2, P, n_cols=4, prec=5)

```

Name: JERUSHLIN JOSE JB
Register Number: 21222240039

State-value function:

00	0.78049	01	0.65854	02	0.53659	03	0.26829
04	0.78049			06	0.41463		
08	0.78049	09	0.78049	10	0.70732		
		13	0.85366	14	0.92683		

```

1 if(np.sum(V1>=V2)==16):
2     print("The Adversarial policy is the better policy")
3 elif(np.sum(V2>=V1)==16):
4     print("The Improved policy is the better policy")
5 else:
6     print("Both policies have their merits.")

```

The Improved policy is the better policy

```

1 def policy_iteration(P, gamma=1.0, theta=1e-10):
2     random_actions=np.random.choice(tuple(P[0].keys()),len(P))
3     pi = lambda s: {s:a for s, a in enumerate(random_actions)}[s]
4     while True:
5         old_pi = {s:pi(s) for s in range(len(P))}
6         V = policy_evaluation(pi, P,gamma,theta)
7         pi = policy_improvement(V,P,gamma)
8         if old_pi == {s:pi(s) for s in range(len(P))}:
9             break
10    return V, pi

```

```

1 optimal_V, optimal_pi = policy_iteration(P)
2

```

```

1 print("Name: JERUSHLIN JOSE JB ")
2 print("Register Number: 21222240039 ")
3 print('Optimal policy and state-value function (PI):')
4 print_policy(optimal_pi, P, action_symbols=('<', '>', 'v', '^'), n_cols=4)

```

Name: JERUSHLIN JOSE JB
Register Number: 21222240039

Optimal policy and state-value function (PI):

Policy:

00	<	01	^	02	^	03	^
04	<			06	<		
08	^	09	>	10	<		
		13	v	14	>		

```

1 print('Reaches goal {:.2f}%. Obtains an average undiscounted return of {:.4f}.'.format(
2     probability_success(env, optimal_pi, goal_state=goal_state)*100,
3     mean_return(env, optimal_pi)))

```

Reaches goal 69.00%. Obtains an average undiscounted return of 0.6900.

```

1 print("Name: JERUSHLIN JOSE JB ")
2 print("Register Number: 21222240039 ")
3 print_state_value_function(optimal_V, P, n_cols=4, prec=5)

```

```
Name: JERUSHLIN JOSE JB
Register Number: 212222240039
State-value function:
| 00 0.82353 | 01 0.82353 | 02 0.82353 | 03 0.82353 |
| 04 0.82353 |           | 06 0.52941 |           |
| 08 0.82353 | 09 0.82353 | 10 0.76471 |           |
|           | 13 0.88235 | 14 0.94118 |           |
```

1 Start coding or [generate](#) with AI.